

Comparing Machine Learning Models for Monthly Stock Return Prediction

Course Project Report

Momen Ismail

Machine Learning for Finance, University of Bonn

Instructor: Prof. Ivan Gufler

July 27, 2026

Abstract

This project compares six models for predicting next-month stock excess returns: a historical mean benchmark, OLS-3, Partial Least Squares, Elastic Net, Random Forest, and Gradient Boosting. I use a monthly panel with 211,059 stock-month observations, 656 tickers, and 484 predictors from February 1990 to December 2025. The model choices are made using chronological training and validation periods, and the final specifications are evaluated once on the 2020–2025 test sample. Random Forest performs best and reduces monthly squared forecast error by about 1.14% relative to the historical mean. OLS-3 provides a very small improvement, Elastic Net reduces all coefficients to zero, and PLS and Gradient Boosting perform worse than the benchmark. The main conclusion is that machine learning can provide a modest improvement, but the result is not stable across years and does not by itself imply a profitable trading strategy.

Keywords: stock return prediction, machine learning, out-of-sample evaluation, Random Forest, asset pricing

1 Introduction

Predicting stock returns is a difficult problem. Monthly returns are noisy, market conditions change, and relationships that appear useful in one period may disappear in another. At the same time, investors and researchers have a large amount of information available, including price history, accounting variables, market conditions, and macroeconomic indicators. This creates a natural setting for comparing traditional statistical models with machine learning methods.

Machine learning is useful in empirical asset pricing for three main reasons. First, expected-return estimation is a prediction problem: information known at time t is used to predict returns at time $t + 1$. Second, the number of possible predictors is large and many variables are correlated. Third, the relationship between characteristics and future returns may be nonlinear and may depend on interactions between firm-level and aggregate variables. These points motivate regularization, dimension reduction, and tree-based methods (Gu et al. 2020).

The research question of this project is:

Can machine learning models predict next-month stock excess returns more accurately than a simple historical-mean benchmark?

I compare six models under the same data and sample design. The historical mean is the benchmark. OLS-3 is a simple linear model using size, book-to-market, and 12-month momentum. PLS and

Elastic Net use all 484 predictors but control the high dimensionality in different ways. Random Forest and Gradient Boosting allow nonlinear relationships and interactions.

The comparison is designed to be realistic for financial prediction. I do not use random train-test splits. Hyperparameters are selected using expanding windows inside the training period, fixed and optimized versions are compared on a later validation period, and the final models are evaluated once on the 2020–2025 test period. This follows the main lesson from the forecasting literature: a model should be judged on genuinely future observations and against a clear benchmark (Welch and Goyal 2008).

The project is mainly a comparison exercise. I do not propose a new algorithm, and I do not claim that the best model creates a profitable trading strategy. The goal is to see which model generalizes best, where the models fail, and whether greater complexity is actually useful in this data set.

The rest of the paper is organized as follows. [Section 2](#) describes the data and preprocessing. [Section 3](#) explains the models and the comparison design. [Section 4](#) presents the validation and test results. [Section 5](#) discusses the main limitations and what I would do differently. [Section 6](#) concludes.

2 Data

2.1 Sources and observation unit

The observation unit is a stock-month. The stock universe is based on a locked S&P 500-style list that contains current and historical tickers available in the project inputs. This is not a perfect point-in-time index-constituent data set, but the same locked universe is used throughout the pipeline.

Daily stock prices and trading volume come from Yahoo Finance. Daily S&P 500 and VIX data are used for market and volatility variables. Monthly Fama–French data provide the risk-free rate. Annual Compustat data provide accounting variables and industry classifications. Eight aggregate predictor states are taken from the Welch–Goyal data set.

Daily histories begin in 1987 so that longer-window variables can be calculated before the main sample starts. The final model-ready panel covers February 1990 to December 2025.

2.2 Prediction target

For stock i in month t , the target is the next-month excess return:

$$y_{i,t+1} = r_{i,t+1} - RF_{t+1}, \quad (1)$$

where $r_{i,t+1}$ is the stock return in the next calendar month and RF_{t+1} is the corresponding risk-free rate.

The target is created only when the next observation for the same ticker is the immediately following calendar month. This avoids treating a return over a gap of several months as a one-month return. The target itself is kept in its raw return units and is not winsorized.

2.3 Timing and leakage controls

The main timing rule is simple: a row dated month t may only use information available by the end of month t .

Stock-price characteristics use price history through month t . Compustat requires an additional timing rule because the fiscal year-end date is not the same as the public reporting date. I therefore treat annual accounting data as available six months after fiscal year-end and use the latest report available for each stock-month. Very old accounting observations are not carried forward indefinitely.

The Welch–Goyal variables are lagged by one month. Therefore, a stock-month in month t uses the aggregate state from month $t - 1$. Together with the exact next-month target rule, these choices reduce look-ahead bias.

2.4 Preprocessing

Missing values in the 45 continuous firm characteristics are replaced with the cross-sectional median of the same month. This uses information from other firms in that month, but not from future months. The continuous characteristics are then winsorized within each month at the 1st and 99th percentiles. This limits the influence of extreme predictor values while keeping the target unchanged.

The three binary characteristics remain binary. Market and VIX variables, Welch–Goyal variables, and SIC2 indicators are not cross-sectionally winsorized. The 360 interaction variables are formed from the processed continuous characteristics and the eight aggregate states.

2.5 Predictors and final sample

The final data contain 484 predictors:

- 45 continuous stock and accounting characteristics;
- 3 binary stock characteristics;
- 4 market and VIX variables;
- 8 lagged Welch–Goyal variables;
- 64 SIC2 industry indicators;
- 360 characteristic–macro interactions.

The continuous characteristics cover size, valuation, profitability, investment, leverage, liquidity, momentum, reversal, volatility, beta, and trading activity. The interactions allow the effect of a firm characteristic to depend on the aggregate economic state.

The final panel contains 211,059 observations and 656 tickers. It has no missing predictor values, infinite values, or duplicated ticker-month observations.

Table 1: Chronological sample design

Sample	Observations	Months	Tickers	Start	End
Training	131,061	299	572	1990-02	2014-12
Validation	35,386	60	611	2015-01	2019-12
Development	166,447	359	613	1990-02	2019-12
Test	44,612	72	641	2020-01	2025-12

The training sample is used for tuning. The validation sample is used to choose between fixed and optimized versions. The development sample combines training and validation for the final fit. The test sample is used only for the final comparison.

3 Methodology

3.1 Prediction framework

The purpose of the models is to use information known at the end of month t to predict the excess return realized in month $t + 1$. I write the forecasting problem as

$$r_{i,t+1}^e = \mathbb{E}_t \left[r_{i,t+1}^e \right] + \varepsilon_{i,t+1}, \quad \mathbb{E}_t \left[r_{i,t+1}^e \right] = g(\mathbf{x}_{i,t}), \quad (2)$$

where $r_{i,t+1}^e$ is the next-month excess return of stock i , $\mathbf{x}_{i,t}$ is the information available at month t , and $\varepsilon_{i,t+1}$ is the part of the return that is not predicted by the model.

The models differ in how they estimate the function $g(\cdot)$. Some impose a simple linear relationship, while others allow regularization, dimension reduction, nonlinearities, and interactions. I evaluate all models using the same target, chronological samples, and performance measures so that their results are directly comparable.

3.2 Model choice and rationale

I chose six models that represent different levels and types of model complexity. The purpose is not only to find the best-performing model, but also to test whether additional complexity is actually useful for this prediction problem.

Table 2: Models included in the comparison

Model	Reason for inclusion
Historical Mean	A simple constant benchmark. It shows whether the other models provide any improvement over the average return observed in the available estimation sample.
OLS-3	A small and interpretable linear model using size, book-to-market, and 12-month momentum. It provides a comparison between a simple economic model and the high-dimensional methods.
Partial Least Squares	A dimension-reduction model that replaces the 484 correlated predictors with a small number of supervised components.
Elastic Net	A regularized linear model that can shrink unstable coefficients and set some of them exactly to zero.
Random Forest	A tree ensemble that can capture nonlinear relationships and interactions while reducing the instability of an individual tree through averaging.
Gradient Boosting	A second tree ensemble that fits shallow trees sequentially, with each new tree trying to correct errors made by the previous trees.

Historical Mean

The historical mean is the main benchmark. It ignores all stock characteristics and assigns the same prediction to every stock. During each training or validation exercise, the mean is calculated using only the permitted estimation sample. For the final test, it is calculated using the complete February 1990–December 2019 development sample.

This is an important benchmark because a machine-learning model is only useful for this project if it improves on a simple constant forecast.

OLS-3

OLS-3 uses three predictors:

- firm size;
- book-to-market;
- 12-month momentum.

I chose these variables because they are well-known and economically interpretable stock characteristics (Fama and French 1992; Jegadeesh and Titman 1993). The model provides a useful middle point between the historical mean and the 484-predictor models.

It also allows me to examine whether a small set of established predictors can perform as well as, or better than, much more complex methods.

Partial Least Squares

Partial Least Squares (PLS) is included because the predictor set is high-dimensional and strongly correlated. PLS constructs a small number of latent components using information from both the predictor matrix and the target. It therefore differs from an unsupervised method such as ordinary Principal Component Analysis, which constructs components without using the target variable (Wold et al. 1984).

The main hyperparameter is the number of PLS components. Too few components may remove useful information, while too many components may reproduce noise from the original predictor set.

Elastic Net

Elastic Net is included as a high-dimensional linear model. It combines two types of regularization:

- the L_1 penalty, which can set coefficients to zero;
- the L_2 penalty, which stabilizes coefficients when predictors are strongly correlated.

The parameter α controls the overall strength of regularization, while the l_1 -ratio determines the balance between the two penalties (Tibshirani 1996; Zou and Hastie 2005).

This model is useful because it tests whether the predictive information can be represented through a sparse or heavily regularized linear relationship.

Random Forest

Random Forest fits many regression trees to bootstrap samples and averages their predictions (Breiman 2001). Each split considers only a random subset of predictors.

The random sampling of both observations and predictors reduces the correlation between the trees. Averaging then reduces the variance of an individual tree. The model can still represent nonlinearities and interactions without requiring these relationships to be specified in advance.

The Random Forest specification also uses out-of-bag observations. I use the out-of-bag result only as an additional development-sample diagnostic. It is not treated as final test evidence and is not used instead of the chronological validation and test samples.

Gradient Boosting

Gradient Boosting also combines regression trees, but it does so sequentially. Each new tree is fitted to improve the errors remaining from the existing ensemble.

The main decisions are:

- the number of trees;
- the learning rate;
- the maximum depth of each tree.

A small learning rate makes each tree contribute only a limited correction. Shallow trees restrict the complexity of each individual learner. The number of trees determines how many sequential corrections the model is allowed to make.

3.3 Model-specific scaling and implementation choices

The continuous firm characteristics were already median-imputed and winsorized during data construction. For model estimation, I make an additional scaling decision.

The predictors used by OLS-3, PLS, and Elastic Net are standardized using the mean and standard deviation of the relevant estimation sample. The same fitted transformation is then applied to the corresponding validation or test sample.

Standardization is important for PLS and Elastic Net because their components and penalties depend on the scale of the predictors. I also standardize OLS-3 for consistency, although standardization does not change the fitted values of an unregularized OLS regression with an intercept.

Random Forest and Gradient Boosting are not standardized. Tree splits depend on the ordering of predictor values rather than their measurement units, so rescaling the variables is unnecessary.

Both tree models use a fixed random seed of 42 so that the results can be reproduced.

3.4 Tuning design

I do not use random cross-validation. A random split would mix earlier and later market observations and could allow future market conditions to influence model selection.

Instead, hyperparameters are selected using ten annual expanding-window folds inside the February 1990–December 2014 training sample:

- the first fold trains through December 2004 and validates on 2005;
- the training window then expands by one year for each new fold;
- the final fold trains through December 2013 and validates on 2014.

For each candidate specification, I calculate monthly MSE in every validation year and then average the ten annual results:

$$\text{CV-MSE}(\lambda) = \frac{1}{10} \sum_{k=1}^{10} \text{MSE}_{\text{month},k}(\lambda), \quad (3)$$

where λ represents the hyperparameter combination.

The tuning stage examines:

- the number of components for PLS;
- α and the l_1 -ratio for Elastic Net;
- the number of trees for Random Forest;
- the number of trees, learning rate, and maximum depth for Gradient Boosting.

I use monthly MSE as the tuning criterion because the number of stocks differs across months. Averaging errors within month first prevents months with more observations from receiving greater weight.

3.5 Fixed and optimized specifications

Before tuning, I estimate fixed reference versions of the models. These provide a direct test of whether optimization actually improves performance.

The main fixed specifications are:

- PLS with 20 components;
- Elastic Net with $\alpha = 0.0001$ and l_1 -ratio = 0.50;
- Random Forest with 100 trees;
- Gradient Boosting with 100 trees, learning rate 0.01, and maximum depth 2.

The expanding-window training exercise produces optimized candidates. However, I do not automatically send these candidates to the final test. I first compare each optimized specification with its fixed version using the later January 2015–December 2019 validation sample.

This additional validation step is important because a specification that performs well in earlier training folds may not remain useful in a later market period.

3.6 Validation-based final model selection

The validation results are shown in [Table 3](#).

Table 3: Validation choice between fixed and optimized specifications

Model family	Fixed monthly MSE	Optimized monthly MSE	Final validation choice
PLS	0.0077	0.0076	Optimized version with 2 components
Elastic Net	0.0078	0.0074	Optimized version with $\alpha = 0.015$ and l_1 -ratio = 0.85
Random Forest	0.0075	0.0075	Fixed version with 100 trees
Gradient Boosting	0.0075	0.0076	Fixed version with 100 trees, learning rate 0.01, and depth 2

The results lead to four different decisions.

First, the optimized PLS model is retained because two components produce a lower validation MSE than the fixed 20-component model. This suggests that a small number of supervised components generalizes better than a larger component representation.

Second, the optimized Elastic Net is retained because it improves on the fixed Elastic Net in validation. The selected penalty is strong enough to shrink all 484 coefficients to zero. I do not

treat this as a coding failure. It means that, under the selected chronological validation design, the model preferred a constant prediction over unstable linear predictor effects.

Third, expanding-window tuning favored a larger Random Forest candidate with 300 trees, but the fixed 100-tree model produces a lower MSE in the separate 2015–2019 validation period. I therefore retain the fixed 100-tree version.

Fourth, the more complex 300-tree Gradient Boosting candidate also performs worse than the fixed 100-tree model in the later validation period. I therefore retain the fixed version.

These decisions show why I use both tuning folds and a separate validation sample. Optimization inside the training period does not automatically imply better performance in a genuinely later period.

The final Random Forest specification is:

- 100 trees;
- square-root feature sampling at each split;
- minimum leaf size of 20;
- bootstrap sampling;
- out-of-bag diagnostics;
- random seed 42.

The final Gradient Boosting specification is:

- 100 trees;
- learning rate 0.01;
- maximum tree depth 2;
- random seed 42.

All final choices are made using the training and validation periods. The 2020–2025 test results are not used to choose these specifications.

3.7 Final estimation and test procedure

After the final specifications are selected, I combine the training and validation samples into the February 1990–December 2019 development sample.

The final procedure is:

1. select hyperparameters using the 1990–2014 training sample;
2. compare fixed and optimized models using the 2015–2019 validation sample;
3. lock the final specifications;
4. refit each selected model on the complete 1990–2019 development sample;
5. evaluate each fitted model once on the January 2020–December 2025 test sample.

The final test contains 44,612 stock-month observations over 72 months. After the test results are opened, I do not change the predictors, hyperparameters, sample dates, or model specifications.

3.8 Evaluation metrics

The primary evaluation measure is monthly mean squared error:

$$\text{MSE}_{\text{month}} = \frac{1}{T} \sum_{t=1}^T \left[\frac{1}{N_t} \sum_{i=1}^{N_t} (y_{i,t+1} - \hat{y}_{i,t+1})^2 \right]. \quad (4)$$

The squared errors are first averaged across stocks within each month and then across months. This gives every month the same weight, even when the number of stocks differs.

The main benchmark-relative measure is monthly out-of-sample R^2 :

$$R_{\text{OOS,month}}^2 = 1 - \frac{\text{MSE}_{\text{month,model}}}{\text{MSE}_{\text{month,HM}}}. \quad (5)$$

A positive value means that the model reduces average monthly squared forecast error relative to the historical mean. A negative value means that it performs worse than the benchmark.

The benchmark in this project is the historical mean estimated from the permitted estimation sample. It is not a zero-return forecast. Therefore, the out-of-sample R^2 measures improvement relative to an estimated constant expected-return forecast.

I also report:

- monthly RMSE;
- pooled RMSE;
- pooled MAE;
- pooled out-of-sample R^2 ;
- prediction-target correlation.

The monthly measures give equal weight to months, while the pooled measures give equal weight to individual stock-month observations. I treat the monthly measure as the main criterion and the pooled measures as supporting evidence.

Correlation is undefined for constant predictions. This applies to the historical mean and to the final Elastic Net model, which produces an effectively constant prediction.

3.9 Scope of the comparison

The model ranking is descriptive. I do not conduct a formal statistical test showing that the leading model is significantly better than every competing model.

I also do not construct investment portfolios or calculate turnover, transaction costs, Sharpe ratios, or cumulative trading returns. The methodology therefore compares statistical forecast accuracy, not investment profitability.

Finally, I do not create an additional ensemble that combines the predictions of all six models. Random Forest and Gradient Boosting are themselves ensemble methods, but the final comparison evaluates each of the six approaches separately.

4 Results

4.1 Overall final-test performance

After completing the tuning and validation stages, I locked the final model specifications and refitted them on the complete February 1990–December 2019 development sample. I then evaluated them once on the January 2020–December 2025 test sample.

Table 4 reports the final results. The models are ranked by monthly MSE, which is the main evaluation criterion. The two out-of-sample R^2 columns are reported in percent to make the relatively small differences easier to read.

Table 4: Final test performance, 2020–2025

Rank	Model	Monthly MSE	Monthly RMSE	Monthly OOS R^2 (%)	Pooled OOS R^2 (%)	Correlation
1	Random Forest	0.022112	0.148700	1.1427	1.1071	0.1138
2	OLS-3	0.022360	0.149534	0.0301	0.0296	0.0313
3	Elastic Net	0.022367	0.149556	0.0000	0.0000	–
4	Historical Mean	0.022367	0.149556	0.0000	0.0000	–
5	PLS	0.022509	0.150032	-0.6363	-0.6579	0.0270
6	Gradient Boosting	0.022573	0.150242	-0.9192	-0.9203	-0.0481

Notes: Models are ranked by monthly MSE. The reported out-of-sample R^2 values have been multiplied by 100 and are therefore shown as percentages. Monthly OOS R^2 compares monthly MSE with the historical-mean benchmark. Pooled OOS R^2 uses all stock-month squared errors. Correlation is undefined for constant predictions.

Random Forest ranks first according to the prespecified monthly MSE criterion. Its monthly out-of-sample R^2 is 1.1427%. This means that the model reduces average monthly squared prediction error by approximately 1.14% relative to the historical mean. It does not mean that the model predicts a return of 1.14%.

The pooled out-of-sample R^2 is similar at 1.1071%. The monthly measure gives every month equal weight, while the pooled measure gives every stock-month observation equal weight. The similar values show that the result is not driven only by months containing more stocks.

OLS-3 ranks second, but its improvement is very small. Its monthly out-of-sample R^2 is only 0.0301%. Nevertheless, this is an important comparison because OLS-3 uses only size, book-to-market and 12-month momentum. A simple three-variable model therefore remains competitive with the much larger models.

Elastic Net is effectively tied with the historical mean. The selected penalty shrinks all 484 coefficients to zero, so the model produces a constant prediction. I interpret this as a valid negative result rather than a coding error. Under the selected penalty and chronological validation design, Elastic Net did not find stable linear effects that improved on the benchmark.

PLS and Gradient Boosting both have negative out-of-sample R^2 . Their predictions therefore produce greater squared error than the historical mean. Gradient Boosting also has a negative prediction-target correlation. These results show that additional model complexity does not automatically improve out-of-sample performance.




Figure 1: Monthly out-of-sample R^2 relative to the historical-mean benchmark in the 2020–2025 test sample. The graph displays the R^2 values in decimal form, while Table 4 reports them in percent. Random Forest’s value of approximately 0.011427 corresponds to an error reduction of approximately 1.1427%.

Figure 1 makes the benchmark comparison visible. Random Forest has the clearest positive result. OLS-3 is only slightly above zero, Elastic Net overlaps with the historical mean, and PLS and Gradient Boosting are below zero.

The ranking is descriptive. I do not estimate confidence intervals for the differences between models and I do not conduct a multiple-model forecast comparison test. I can therefore say that Random Forest ranks first according to monthly MSE, but I cannot claim that it is statistically significantly better than every alternative.

4.2 Performance across test years

The complete test-period result does not mean that Random Forest performs well in every year. Figure 2 reports the monthly out-of-sample R^2 separately for each of the six test years.




Figure 2: Year-specific monthly out-of-sample R^2 for 2020–2025. Random Forest ranks first over the complete test period but has a positive annual value in only one of the six years. The historical mean and the near-constant Elastic Net predictions overlap close to zero.

Random Forest has a positive annual out-of-sample R^2 only in 2020. Its result in that year is sufficiently strong to offset several negative annual values and produce a positive result over the complete test period.

OLS-3 has a much smaller full-period improvement, but its annual out-of-sample R^2 is positive in four of the six test years. PLS and Gradient Boosting are positive in only one year each.

This comparison changes how I interpret the overall ranking. Random Forest has the lowest average forecast error over the complete period, but its advantage is concentrated in one year and is not stable across all market environments. OLS-3 produces a much smaller average gain, but its yearly performance is more frequently above the benchmark.

The yearly results therefore support a cautious conclusion. Random Forest performs best according to the main full-period measure, but the evidence does not show that it will dominate the simpler models in every future period.

4.3 What the Random Forest uses

Because Random Forest is the highest-ranked model, I also examine where its fitted prediction variation comes from. Individual predictor names are often technical, especially for the interaction variables, so I first summarize importance by broad predictor group.




Figure 3: Random Forest impurity importance aggregated by predictor group in the final development-sample fit. The values describe the fitted model and should not be interpreted as causal effects or estimated risk premia.

Characteristic–macro interactions account for approximately 73.7% of the total Random Forest impurity importance. The Welch–Goyal variables account for approximately 11.7%, market and VIX variables for 7.8%, and the remaining stock characteristics for approximately 6.8%. SIC2 industry indicators contribute almost no measured importance.

This pattern suggests that Random Forest obtains much of its fitted prediction variation from relationships that change with aggregate economic conditions. In other words, the model appears to use combinations of firm characteristics and macroeconomic states rather than relying on one single characteristic.

However, this conclusion has to be treated carefully. The interaction group contains 360 predictors, while the other groups contain far fewer variables. Its high total importance therefore does not mean that every interaction is individually important.

Impurity importance also has general limitations. Correlated variables can divide importance among themselves, and the values can depend on the number of possible split points. The figure describes how the fitted Random Forest makes predictions; it does not identify causal effects.

4.4 Comparison with the machine-learning asset-pricing literature

The general motivation of this project follows the machine-learning asset-pricing literature. Flexible models may be useful because they can handle many predictors and allow nonlinear relationships and interactions (Gu et al. 2020).

The results of this project are mixed. Random Forest improves on the historical-mean benchmark, while Gradient Boosting and PLS do not. OLS-3 also remains competitive despite using only three predictors. The results therefore support the idea that flexible models can help, but they also show that greater complexity is not always rewarded outside the estimation sample.

The numerical results should not be directly compared with percentages reported in another study. The stock universe, sample period, predictors, preprocessing, and benchmark can differ. In particular, the out-of-sample R^2 in this project uses the estimated historical mean as its benchmark,

rather than a zero-return benchmark.

4.5 Financial interpretation

The project compares statistical forecast accuracy. It does not evaluate a trading strategy.

A higher out-of-sample R^2 does not automatically imply a higher Sharpe ratio or a profitable long-short portfolio. A model may reduce average squared prediction error without producing a useful ranking after trading costs and portfolio constraints are considered.

The current analysis does not include:

- portfolio formation;
- transaction costs;
- portfolio turnover;
- liquidity or market-impact costs;
- short-selling restrictions;
- risk-adjusted portfolio returns;
- Sharpe ratios or cumulative trading returns.

The financial conclusion is therefore limited. Random Forest has the lowest forecast error in the 2020–2025 test period, but the project does not show that trading on its predictions would be profitable.

A portfolio analysis would be a separate extension. It would require a predefined sorting or weighting rule and an evaluation of both gross and net returns. It should not be designed after examining the same test results, because that would use the test sample for a second model-selection exercise.

5 Critical Reflection

5.1 What worked

The chronological design is one of the strongest parts of the project. The models are tuned using expanding windows, fixed and optimized versions are compared on a later validation period, and the final test period is kept separate from these choices. This is more realistic than a random split for financial time-series data.

The comparison also produces useful negative results. Elastic Net selects no variables, PLS underperforms, and Gradient Boosting performs worse than the benchmark. These outcomes show that more predictors and more flexible models do not automatically improve forecasts.

Random Forest is the only model with a clear positive full-period result. Its ability to use nonlinearities and interactions appears useful, especially when firm characteristics are combined with aggregate states.

5.2 Main limitations

The first limitation is the stock universe. The project uses a locked S&P 500-style ticker list, not a fully point-in-time CRSP constituent universe. Ticker changes, mergers, delistings, and historical membership may therefore create survivorship or selection bias.

Second, Yahoo Finance is convenient but does not provide the security master, delisting returns, and microstructure detail available in CRSP. Adjusted prices can also be revised. Compustat is lagged by six months, but exact filing dates are not used.

Third, the target has large outliers. The maximum monthly excess return is much larger than a typical observation, so MSE can be strongly affected by rare returns. Reporting MAE helps, but the main ranking still depends on squared loss.

Fourth, the test period contains only six years and includes unusual market conditions. The yearly figure shows that the Random Forest improvement is not stable across the full period.

Fifth, the project does not test statistical significance of loss differences. Random Forest ranks first, but the difference from OLS-3 may not be statistically significant after accounting for time dependence and comparison across several models.

Finally, the project does not evaluate cross-sectional portfolio performance. A useful additional metric would remove each month's average return and test whether the model predicts which stocks perform relatively better or worse within the month. A separate portfolio analysis could then sort stocks by predicted returns and report returns, Sharpe ratios, turnover, and transaction costs.

5.3 What I would do differently

With more time and better data access, I would make five main improvements:

- use a point-in-time CRSP/Compustat universe with permanent identifiers and delisting returns;
- use exact public filing dates instead of a fixed six-month accounting lag;
- keep a longer untouched test period or evaluate several historical forecast origins;
- add formal tests for differences in monthly prediction loss;
- design a portfolio exercise before opening a new test sample and include realistic trading costs.

I would also compare predictor groups more systematically. Instead of relying only on impurity importance from one fitted Random Forest, I would examine how importance changes across time windows and whether the same groups remain useful in different market regimes.

6 Conclusion

This project compares six models for predicting next-month stock excess returns. The main question is whether machine learning improves on a simple historical-mean forecast.

The answer is cautious. Random Forest performs best in the 2020–2025 test sample and reduces average monthly squared forecast error by about 1.14%. OLS-3 provides only a very small improvement but remains competitive despite using only three predictors. Elastic Net shrinks all coefficients to zero and matches the benchmark. PLS and Gradient Boosting perform worse than the historical mean.

The comparison shows that complexity is not automatically valuable. Random Forest benefits from averaging many trees and using interactions, but Gradient Boosting and PLS do not generalize well in this sample. The yearly results also show that Random Forest's full-period advantage is not stable across all six years.

The main conclusion is therefore not that machine learning solves stock return prediction. It is that one ensemble method produces a modest improvement under a careful chronological design,

while several other sophisticated methods do not. Simple benchmarks, validation on future periods, and honest reporting of failures are essential when comparing financial prediction models.

References

- Breiman, L. (2001). “Random Forests”. In: *Machine Learning* 45.1, pp. 5–32.
- Fama, E. F. and K. R. French (1992). “The Cross-Section of Expected Stock Returns”. In: *The Journal of Finance* 47.2, pp. 427–465.
- Gu, S., B. Kelly, and D. Xiu (2020). “Empirical Asset Pricing via Machine Learning”. In: *The Review of Financial Studies* 33.5, pp. 2223–2273.
- Jegadeesh, N. and S. Titman (1993). “Returns to Buying Winners and Selling Losers: Implications for Stock Market Efficiency”. In: *The Journal of Finance* 48.1, pp. 65–91.
- Tibshirani, R. (1996). “Regression Shrinkage and Selection via the Lasso”. In: *Journal of the Royal Statistical Society: Series B* 58.1, pp. 267–288.
- Welch, I. and A. Goyal (2008). “A Comprehensive Look at the Empirical Performance of Equity Premium Prediction”. In: *The Review of Financial Studies* 21.4, pp. 1455–1508.
- Wold, S. et al. (1984). “The Collinearity Problem in Linear Regression: The Partial Least Squares Approach to Generalized Inverses”. In: *SIAM Journal on Scientific and Statistical Computing* 5.3, pp. 735–743.
- Zou, H. and T. Hastie (2005). “Regularization and Variable Selection via the Elastic Net”. In: *Journal of the Royal Statistical Society: Series B* 67.2, pp. 301–320.

A Code and Reproducibility Overview

The full code is included in the project repository. The empirical workflow is organized into separate stages so that data construction, model comparison, and reporting can be checked independently.

1. `src/data/`: daily data cleaning, monthly feature construction, accounting and macro merges, interaction construction, imputation, and winsorization;
2. `src/models/step_01_fixed/`: fixed model specifications;
3. `src/models/step_02_tuning/`: expanding-window tuning;
4. `src/models/step_03_optimization/`: validation comparison of fixed and optimized versions;
5. `src/models/step_04_test/`: final fit on the development sample and one-time test evaluation;
6. `src/models/step_05_interpretation/`: tables, feature importance, yearly results, and figures.

The final model-ready data contain 211,059 rows and 484 predictors. The final paper tables are embedded in the LaTeX source, while the three main figures are stored under `documentation/figures/`. This allows the PDF to compile without requiring the large model-output folder.

The official test results should not be used to retune models or change the predictor set. Re-running the same fixed workflow is a reproducibility check, not a new model-selection exercise.